

A STRUCTURED APPROACH TO CONTROLLED INTRODUCTION OF CHANGES

AGENDA

- nomenclature
- phased rollout success condition
 - configuration validity & service health
- life cycle management / introducing change
 - problem definition
- a solution: feature-flags

AUTOMATION CHANGES OVER TIME

- automation is the embodiment of processes in programmatic fashion
- products and processes change, thus automation must change over time

NOMENCLATURE

- blast radius containment
- phased rollout

BLAST RADIUS CONTAINMENT

- strictly about limiting impact of a change
- for example, make sure we cannot commit to more than X devices in one transaction
 - this would prevent making one network wide transaction changing all instances of Y at the same time

PHASED ROLLOUT

- method of rolling out a change in controlled fashion by applying a subset of the change at a time
- for example, total change is to 1000 devices
 - start by changing one device
 - validate change went well
 - proceed with changing next device
 - validate and proceed to next
 - repeat for all devices...
 - or abort on failure

PHASED ROLLOUT SUCCESS CONDITION

- key to phased rollout is concept of **success condition**
- we roll out change in small parts - sliding window
- how do we know if a change went well?
- how do we know if we can proceed and configure next device?

CONFIGURATION VALIDITY

- configuration commit only includes syntax and semantic checks
 - an empty configuration is valid
 - but would lead to unhealthy device / service
- thus, need concept of **success condition** beyond basic config validity

SERVICE HEALTH

- need to understand if service is **healthy**
- monitor operational state of service
 - is BGP neighbor up?
 - is interface up?
 - can we ping?
- service specific! not generic...

INTRODUCING (NAIVE) CHANGE

- change service configuration template
 - git commit
 - deploy new NSO package
 - re-deploy service instances
- > new config is now active in network

EXAMPLE CHANGE

- modify MTU of service from default (1500) to 9100

CHANGE SERVICE TEMPLATE

```
<config-template xmlns="http://tail-f.com/ns/config/1.0">
  <devices xmlns="http://tail-f.com/ns/ncs">

    <device tags="nocreate">
      <name>{/device}</name>
      <config tags="merge">
        <interface-configurations xmlns="http://cisco.com/ns/yang/Cisco-IOS-XR-ifmgr-cfg">
          <interface-configuration>
            <active>act</active>
            <interface-name>{/interface}</interface-name>
            <description>Link to {/remote/device} [{/remote/interface}]</description>
            <mtus>
              <mtu>
                <owner>{$INTERFACE TYPE}</owner>
                <!-- new hard-coded MTU -->
                <mtu>9100</mtu>
              </mtu>
            </mtus>
            <shutdown tags="delete" when="{/shutdown='false'}"/>
            <!-- ... other config stuff ... -->
          </interface-configuration>
        </interface-configurations>
      </config>
    </device>
  </devices>
</config-template>
```

BENEFIT OF HARD-CODED VALUE


- **no choice means avoiding** test cases - good!
- configurable MTU means we need to test all or reasonable set of values
- combinatorial explosion with many config knobs

A close-up photograph of three handguns and several bullets on a rustic wooden surface. In the foreground, a black semi-automatic handgun is positioned diagonally. Behind it, a silver revolver is visible. To the left, another black handgun is partially shown. Several bullets are scattered around the guns. The text "LOADED GUN SCENARIO" is overlaid in white, bold, sans-serif font across the center of the image.

LOADED GUN SCENARIO

LATENT CHANGE = LOADED GUN

- latent change in service code = loaded gun
 - starts when service package is deployed
 - end when service is re-deployed




KALLE FIRED THE GUN

- anyone else coming in doing (trivial) change in loaded gun window will inadvertently push MTU change
 - Kalle wants to fix spelling mistake in description
 - considered trivial, didn't do `commit dry-run`
 - pushes MTU change causing service outage

NO REVERT

- template change & re-deploy moves forward
- no way back
 - except rollback, but only works for naive scenario
 - interleaved transactions make rollback useless



FEATURE GROUPING

GOALS

- no loaded gun
- going backwards rollback
- success condition / service health
- avoid combinatorial explosion
 - allows testing

FEATURE-FLAGS

- well known concept in software development
- move introduction of change from commit/deploy time to run time
 - temporal decoupling of development and operations!!!
- focus on transition / change
- limited life time

FEATURE-FLAG

- emphasize old -> new transition

```
list backbone-interface {  
  key "device interface";  
  // other things  
  
  container feature-flags {  
    leaf high-mtu {  
      type boolean;  
      description "Enable new high MTU (9100). Disable for old MTU (1500)";  
      default "false";  
    }  
  }  
}
```

FEATURE-FLAG IN TEMPLATE

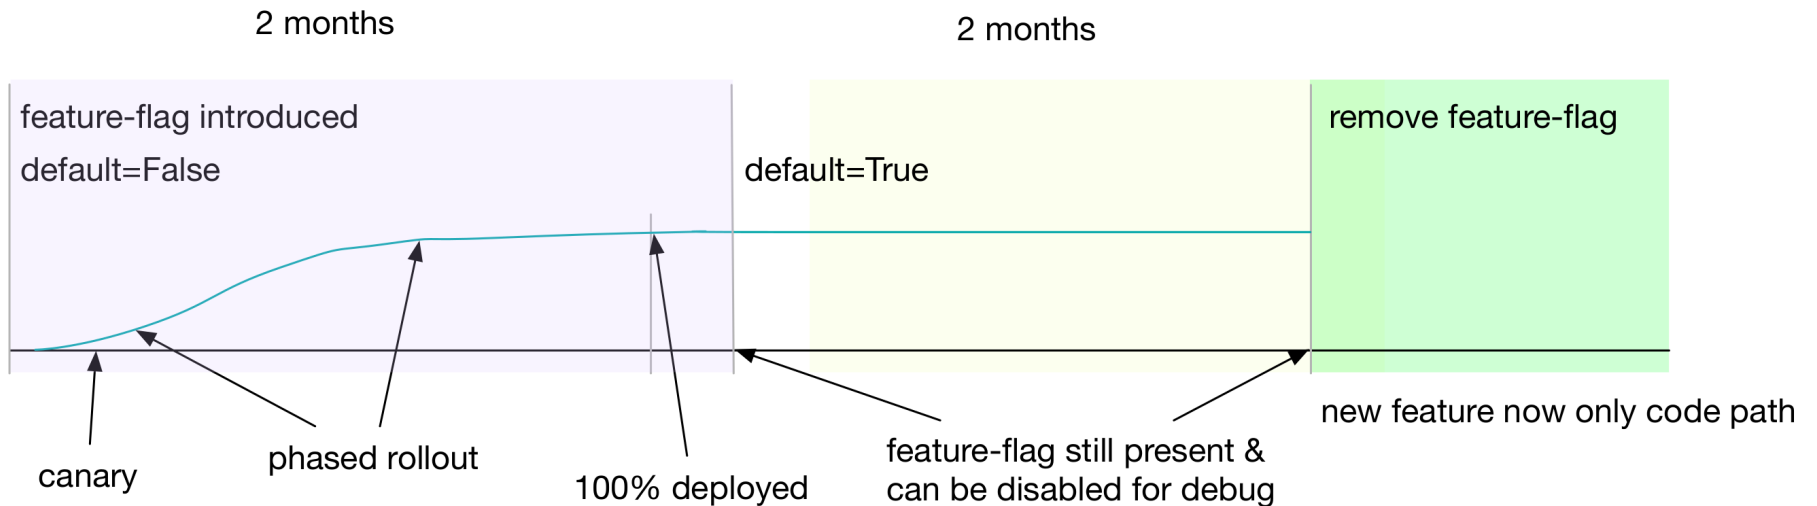
```
<config-template xmlns="http://tail-f.com/ns/config/1.0">
  <devices xmlns="http://tail-f.com/ns/ncs">

    <device tags="nocreate">
      <name>{/device}</name>
      <config tags="merge">
        <interface-configurations xmlns="http://cisco.com/ns/yang/Cisco-IOS-XR-ifmgr-cfg">
          <interface-configuration>
            <active>act</active>
            <interface-name>{/interface}</interface-name>
            <description>Link to {/remote/device} [{/remote/interface}]</description>
            <mtus>
              <mtu>
                <owner>{$INTERFACE TYPE}</owner>
                <!-- new high MTU conditioned on feature-flags -->
                <mtu when="/feature-flags/high-mtu='true'">9100</mtu>
              </mtu>
            </mtus>
            <shutdown tags="delete" when="{/shutdown='false'}"/>
            <!-- ... other config stuff ... -->
          </interface-configuration>
        </interface-configurations>
      </config>
    </device>
  </devices>
</config-template>
```

SOCIOTECHNICAL

- technically, FF is *just another input*
- NSO won't treat it differently
- difference is in concept
 - clear life cycle for FF
 - introduce FF for change transition
 - when done, **remove** FF
 - keeps down input / permutations over time

FEATURE-FLAG LIFE CYCLE



ANTI-PATTERN

- we could introduce new MTU leaf
- allows arbitrary pick of MTU
- BAD - we want choice of 1500 or 9100
- reduce choice / permutations

```
list backbone-interface {  
  key "device interface";  
  // other things  
  
  leaf mtu {  
    type uint16 {  
      range "1500..9100";  
    }  
    description "MTU of service";  
    default "1500";  
  }  
}
```

ANTI-PATTERN

- better, reduction of choice to 2 values
- still, over time, new config knobs leads to combinatorial explosion
- focus on transitional nature

```
list backbone-interface {  
  key "device interface";  
  // other things  
  
  leaf mtu {  
    type uint16 {  
      range "1500 | 9100";  
    }  
    description "MTU of service, either 1500 (old) or 9100 (new)";  
    default "1500";  
  }  
}
```

PHASED ROLLOUT

- feature-flags are per service instance
- many flags to flip for 10000 service instances
- automatic?

SLIDING WINDOW

- sliding window relies on **success condition**
- in practice: service self-test

PROCEDURE

- find list of feature-flags
- for each;
 - go to service owning feature-flag
 - run service self-test, early exit on fail
 - flip feature-flag
 - run service self-test
 - rollback on error (flip back flag)
 - continue to next FF instance on success

INTROSPECTION

- feature-flag can be specific type
 - makes it a generic procedure to find through introspection
 - flag type can indicate *direction*

```
typedef ff-boolean-false-to-true {  
  type boolean;  
  description "A boolean feature flag that transitions from false to true";  
  default false;  
}
```

FEATURE-FLAG NAVIGATOR

- this is a mock-up

```
show feature-flags feature-flags
```

feature-flag	type	progress
/infrastructure/base-config/feature-flags/foobar	false-to-true	73%
/infrastructure/backbone-interface/feature-flags/bar	false-to-true	14%

```
show feature-flags instances
```

instance	type	value	complete
/infrastructure/base-config{901-R1-2053}/feature-flags/foobar	false-to-true	false	false
/infrastructure/base-config{901-R1-2054}/feature-flags/foobar	false-to-true	true	true
/infrastructure/bb-intf{901-R1-2053 et-9/0/0}/feature-flags/bar	true-to-false	false	true
/infrastructure/bb-intf{901-R1-2053 et-10/0/0}/feature-flags/bar	true-to-false	true	false

SERVICE SELF-TEST

- what is it in practice?
- a YANG action!
- specific to service type
- by returning simple common structure it can be used in generic fashion

RETURN GENERIC

```
action self-test {  
  tailf:info "Perform self-test of the service";  
  tailf:actionpoint "backbone-interface-self-test";  
  output {  
    leaf success {  
      type boolean;  
    }  
  
    container interface {  
      // service specific health / state about the interface  
    }  
    container is-is {  
      // service specific health / state about IS-IS  
    }  
    container pim {  
      // service specific health / state about PIM  
    }  
  }  
}
```

```
action self-test {  
  tailf:info "Perform self-test of the service";  
  tailf:actionpoint "ibgp-neighbor-self-test";  
  output {  
    leaf success {  
      type boolean;  
    }  
  
    container bgp {  
      // BGP specific health / state  
    }  
  }  
}
```

EXAMPLE

```
def get_state(kp_unused, log, root=None, service=None, action_output=None):  
    log.info("get_state for {} {}".format(service.device, service.interface))  
    dev = root.devices.device[service.device]  
    os = utils.get_dev_os(dev)  
    state = service.state
```